A Persistence-anchored evaluation

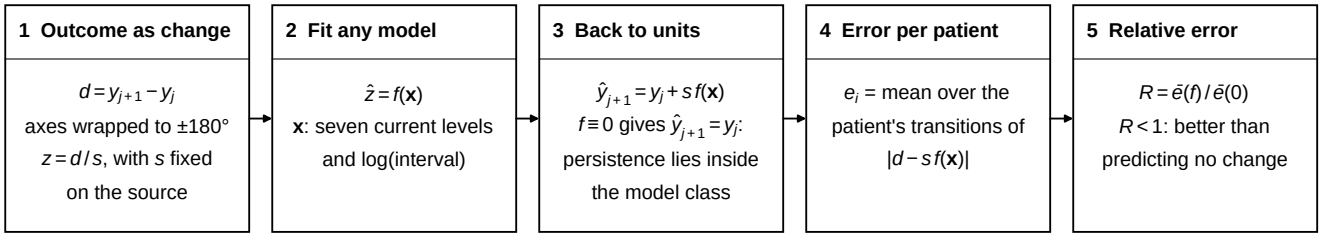

Steps 1 and 3 place the comparator, which needs no data, inside the model class; steps 4 and 5 make the result a statement about patients rather than records. No step depends on the learner.

B Strategies compared at each local cohort size

Strategy	Uses source model	Uses target patients	What is fitted in the target centre	Shrunk toward*
Frozen (reference)	Yes	No	Nothing	–
Recalibration (least squares)	Yes	Yes	One slope per measurement	–
Recalibration (absolute error)	Yes	Yes	One slope, weighted-median fit	–
Recalibration, shrunk	Yes	Yes	One slope	Slope = 1
Local training	No	Yes	Full boosting model	–
Local training, shrunk	No	Yes	Full boosting model	Zero change
Augmented	Yes	Yes	Boosting model, source output as input	–
Residual boosting	Yes	Yes	Boosting model of the residual	–
Residual boosting, shrunk	Yes	Yes	Boosting model of the residual	The frozen model
Pooled source + target	Source data	Yes	Boosting model on both cohorts	–

*Shrinkage weight chosen by five-fold cross-validation within the local cohort, with folds grouped by patient.